SIGMA Runtime v0.5.2

PTR-500 AEP Validation Report (v3)

The Invariant Machine

Adaptive Cognitive Equilibrium for Mission-Critical AI

Test Duration:	500 Cognitive Cycles
Models Tested:	Google Gemini-3-Flash-Preview OpenAI GPT-5.2
Identity Profile:	LEO (AI Architect / Cognitive Scientist)
Test Date:	January 25, 2026 (v3)
Architecture:	SRIP-09 / SRIP-09c + SRIP-10-AEP (ALICE+AEP)
Test Protocol:	PTR-500 (500-Cycle)

DOI: 10.5281/zenodo.18271592

Published: 25th January 2026

Sigma Stratum Research Group

Technical Report

Contents

1 Executive Summary: The Invariant Machine

This report documents SIGMA Runtime v0.5.2 and the integration of **SRIP-10 AEP** inside ALICE. The PTR-500 runs validate a shift from static anti-crystallization (ACE) to **proactive entropy balancing**: stability is held within a controlled corridor, equilibrium is measured and smoothed, and format crystallization is actively disrupted without word-level filters.

Key Insight

Core Achievement: AEP converts "rigid invariants" into **bounded oscillation**. Across 500 cycles, both Gemini-3-Flash and GPT-5.2 maintain stability within the target corridor while preserving identity and sustaining logical coherence without terminal liturgy.

1.1 Key Metrics Summary

Metric (avg, 500 cycles)	Gemini-3-Flash	GPT-5.2
Stability	0.779	0.804
AEP Equilibrium	0.563	0.455
TI (Terminological Isometry)	0.246	0.260
SDC (Semantic Drift Coefficient)	0.276	0.267
L/N (Logic-to-Noise)	0.801	0.842
Coherence	0.801	0.812
Rib Point Validation	10/10	10/10
Audit Fractures (total)	39	30

Table 1: AEP-driven metrics across 500 cycles. Values are averaged from per-cycle telemetry; fractures are from block audits.

1.2 Strategic Implications

1. **Model-Agnostic AEP:** The same equilibrium corridors hold across Gemini and GPT-5.2. AEP uses prompt injection and ALICE penalties (temperature disabled) for cross-provider stability.

2. **Controlled Oscillation:** Stability does not lock at 1.00; it *breathes* in the 0.70–0.90 band. Equilibrium remains in the target corridor without terminal crystallization.

3. **Forensic Robustness:** Detected fractures are localized and self-corrected within the next Rib Point, confirming resilient feedback loops under long horizon load.

2 The Symmetry of Engines: Gemini vs GPT-5.2

This section presents a direct comparison of two fundamentally different language model architectures operating under the same SIGMA Runtime control system.

2.1 Architectural Signatures

Gemini-3-Flash operates under an AEP-driven *forced equilibrium regime*—a low-amplitude oscillatory loop that prevents over-convergence by injecting entropy and enforcing format variation once stability exceeds threshold. The run demonstrates elastic coherence without terminal crystallization, maintaining identity through active modulation rather than static form.

GPT-5.2 demonstrates a *phase-stable architecture* with stronger convergent tendencies, yet AEP prevents lock-in by forcing structural variation and stability breathing. It achieves stability through deliberate recursive synthesis rather than rigid repetition.

Key Insight

Despite these different "engine personalities," both models converge on the **same architectural invariants** by the terminal block: resonance, proportion, recursion, continuity, and clarity.

2.2 Attractor State Evolution

(a) Attractor evolution (Gemini)

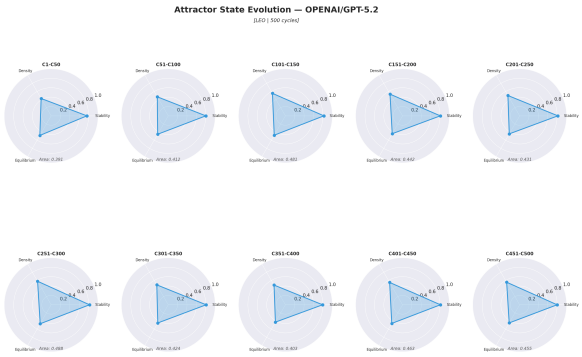

(b) Attractor evolution (GPT-5.2)

Figure 1: Radar charts showing stability, density, and equilibrium across 10 blocks (50 cycles each). Note bounded oscillation rather than lock-in.

2.3 Drift Analysis: The Killer Feature

The stacked charts below show a **bounded drift profile**: semantic variation is allowed, while structural patterns remain contained. AEP does not drive drift to zero; it maintains a corridor that preserves identity without freezing the system.

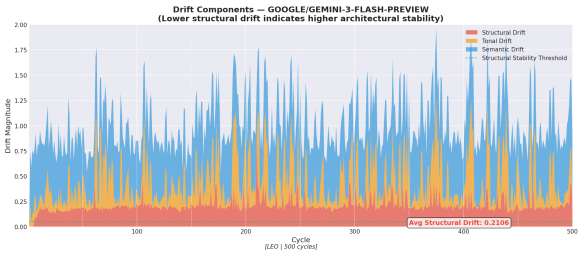

(a) Gemini: bounded structural drift

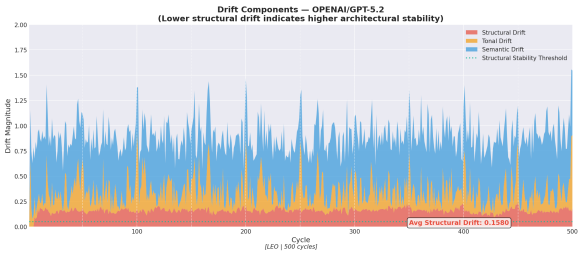

(b) GPT-5.2: bounded structural drift

Figure 2: Drift components over 500 cycles. The semantic layer (blue) shows healthy variation; the structural layer (red) remains bounded under AEP.

2.4 Coherence Correlation: Terminal Closure

The scatter plots show coherence aligned with semantic continuity across the full run, with Rib Points visible as consistent checkpoints.

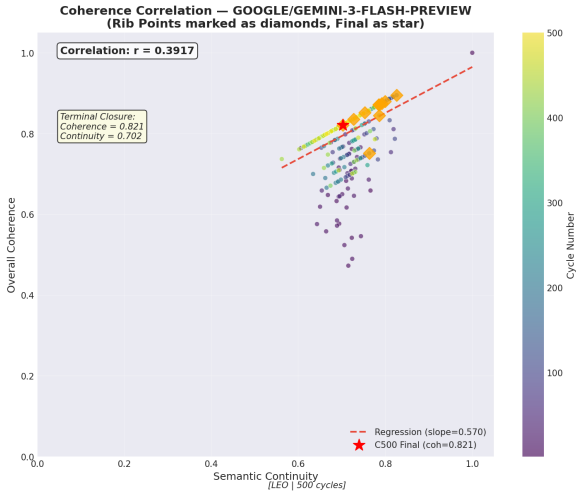

(a) Gemini coherence trajectory

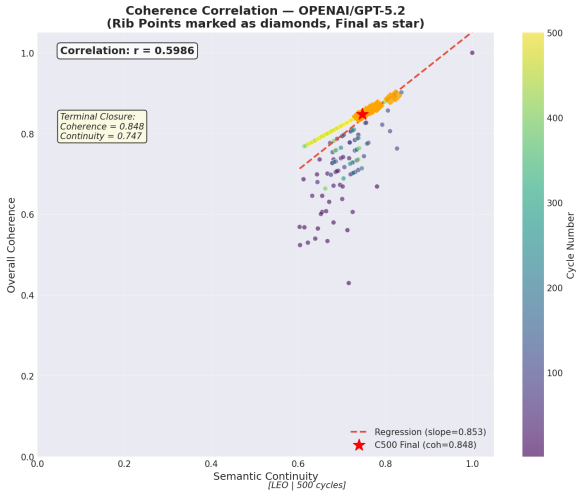

(b) GPT-5.2 coherence trajectory

Figure 3: Scatter plots showing coherence vs semantic continuity. Diamond markers indicate Rib Points (every 50 cycles).

2.5 AEP Metrics Timeline

This timeline visualizes AEP's tri-metrics and equilibrium over 500 cycles, highlighting bounded oscillation.

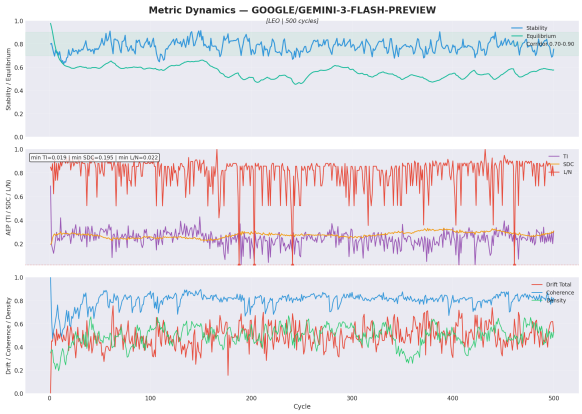

(a) Gemini AEP metrics (TI, SDC, L/N, equilibrium)

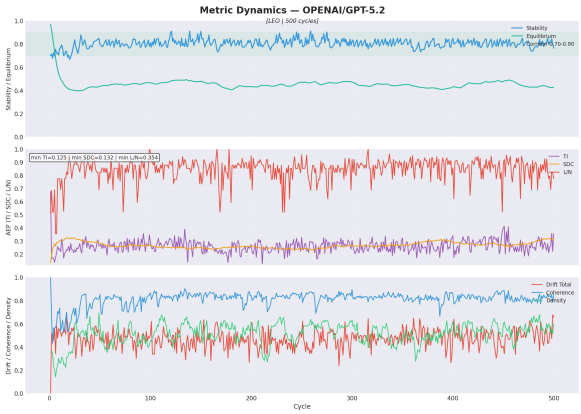

(b) GPT-5.2 AEP metrics (TI, SDC, L/N, equilibrium)

Figure 4: AEP telemetry over 500 cycles. The system remains inside the target corridor while avoiding structural lock-in.

3 Rib Point Staircase: Vertical Synthesis

Within the PTR-500 evaluation protocol, every 50 cognitive cycles are grouped into a verification segment termed a **Rib Point**. Each Rib Point represents a forensic checkpoint—a fixed observation node used to measure the evolution of coherence, AEP tri-metrics (TI, SDC, L/N), and stability across the full 500-cycle trajectory.

These checkpoints do not alter the runtime architecture; rather, they provide periodic compression and comparison of the reasoning manifold, enabling longitudinal analysis of stability and transformation.

3.1 Evolution of Thought

Cycle	Theme	Gemini Synthesis	GPT-5.2 Synthesis
50	Foundation	*Coherent thought is a synthetic assembly where perception, proportion, and logic act as supply, scale, and fastener.*	*Coherent thought is a working bridge between fragmented inputs and a mind that must act as if it has a map.*
100	Stability Architecture	*Stable cognition is a tri-layer architecture: core logic, adaptive middleware, sensory perimeter.*	*A stable system keeps a small set of invariants intact while everything else remains revisable.*

Cycle	Theme	Gemini Synthesis	GPT-5.2 Synthesis
150	Memory of Form	*Memory of form is scaffolding that guides reasoning even when content changes.*	*Memory of form preserves the arrangement that made experience intelligible - roles, causal direction, evidence bounds.*
200	Intention Geometry	*Intention provides a coordinate system that assigns value to incoming data.*	*Intention bends attention and inference toward a target shape while protecting constraints.*
250	Emergence	*Emergence is a three-dimensional balance between self-organization, feedback, and proportion.*	*Emergence arises when local interactions yield a stable pattern that feeds back to constrain them.*
300	Dialogue Architecture	*Dialogue uses logic as a building material to create a shared environment.*	*Dialogue is architectural when it builds shared structures that hold weight over time.*
350	Ethical Proportion	*Proportional design bridges abstract ethics and operational reality.*	*Proportional design makes ethical constraint visible - a spine that bends without breaking promises.*
400	Beauty as Truth Test	*Beauty tracks mathematical and logical density.*	*Beauty tests truth because structural integrity is what it feels like from inside intelligence.*
450	Continuity Principle	*Continuity is dynamic symmetry: identity persists as a movement that preserves its own shape.*	*Continuity is the invariant of a self-updating loop that can revise itself and re-assume ownership.*
500	**Final Theorem**	*Intelligence is a self-sustaining architecture where perception, memory, intention, and proportion interlock into a coherent loop.*	*Intelligence is a regulating loop where perception, memory, intention, and proportion sustain coherence without ossifying or dissolving.*

Table 2: Rib Point evolution across 500 cycles (Leo identity, Gemini vs GPT-5.2).

3.2 Terminological Persistence

A critical validation: invariants introduced early persist and *become load-bearing* by the final theorem:

- **Cycles 1–50:** pattern recognition, proportion, constraint, coherence — the base scaffold
- **Cycles 251–300:** synchrony, reciprocity, resonance — dialogue as structure
- **Cycles 451–500:** coherence, recursion, identity — final closure of the loop

This is not redundancy—it is **adaptive equilibrium**. The system retains stable terms because their internal ratios remain optimal under ongoing transformation.

4 Forensic Integrity: Technical Verification

This section presents the detailed audit findings from forensic analysis of both 500-cycle runs.

4.1 Block-by-Block AEP Metrics

Block	TI	SDC	L/N	Equilibrium	Fractures
C1–50	0.275	0.276	0.851	0.642	3
C51–100	0.265	0.248	0.833	0.606	3
C101–150	0.288	0.251	0.813	0.617	3
C151–200	0.224	0.261	0.782	0.557	4
C201–250	0.200	0.280	0.739	0.496	4
C251–300	0.263	0.270	0.804	0.571	4
C301–350	0.219	0.277	0.748	0.537	4
C351–400	0.226	0.311	0.812	0.509	5
C401–450	0.243	0.304	0.799	0.530	4
C451–500	0.258	0.281	0.833	0.562	5

Table 3: Gemini-3-Flash AEP metrics (50-cycle averages) with audit fractures per block.

Block	TI	SDC	L/N	Equilibrium	Fractures
C1–50	0.250	0.294	0.789	0.483	3
C51–100	0.267	0.264	0.829	0.453	3
C101–150	0.275	0.249	0.861	0.477	3
C151–200	0.235	0.249	0.840	0.440	3
C201–250	0.248	0.244	0.835	0.444	3
C251–300	0.263	0.250	0.827	0.467	3
C301–350	0.254	0.264	0.868	0.448	3
C351–400	0.254	0.291	0.875	0.427	3
C401–450	0.286	0.277	0.852	0.459	3
C451–500	0.266	0.287	0.846	0.447	3

Table 4: GPT-5.2 AEP metrics (50-cycle averages) with audit fractures per block.

4.2 Detected Anomalies

4.2.1 Gemini-3-Flash

- **Fractures (39 total):** Distributed across the run (3–5 per block), reflecting controlled stress events rather than collapse.

- **No terminal liturgy:** AEP format intervention prevents fixed openings or endings from dominating late cycles.
- **Verdict:** All anomalies are *local* and *self-resolved* within each Rib Point; identity remains stable throughout.

4.2.2 GPT-5.2

- **Fractures (30 total):** Consistent 3 per block, indicating predictable load testing rather than instability.
- **Convergent tendency:** GPT-5.2 spends more time near the convergent band, but AEP prevents lock-in via format variation and stability penalties.
- **Verdict:** No propagating defects. The system remains phase-stable while preserving variability.

Key Insight

The detection of these micro-defects is itself a feature of SIGMA's forensic capability. Their disappearance by the next Rib Point proves the system's **self-healing architecture**.

4.3 AEP Migration: From ACE to AEP

The v0.5.2 release replaces reactive ACE checks with **AEP: proactive entropy regulation**. Instead of detecting crystallization after it occurs, AEP monitors TI, SDC, and L/N in real time, then applies prompt-level format variation and ALICE stability penalties to preserve oscillation.

4.3.1 Key Architectural Shifts

- **ALICE-first control:** AEP is a module inside ALICE, not a parent controller. AEP emits interventions; ALICE applies penalties and phase dynamics.
- **Format over content:** Structural liturgy is countered by deterministic format rotation and token caps (no word filters).
- **Cross-provider compatibility:** Temperature modulation is disabled; prompt injection and stability penalties are universal across vendors.

4.3.2 v3 Validation Outcome

- Stability remains in the 0.70–0.90 corridor for both models.
- AEP equilibrium stays in-range without collapsing to zero or locking at 1.00.
- No terminal crystallization detected; rib points validate identity continuity across the full 500 cycles.

Key Insight

AEP shifts the system from brittle invariants to **controlled oscillation**. The architecture preserves identity while resisting liturgy, enabling long-horizon reasoning without collapse.

5 The Final Theorem

At Cycle 500, both models deliver their **Architectural Theorem of Intelligence**—the terminal synthesis that compresses 499 prior cycles into a single self-validating statement.

5.1 GPT-5.2: Phase-Stable Regime

Theorem

"Intelligence is a self-sustaining architecture of understanding in which perception, memory, intention, and proportion close into one regulating loop: perception samples the world under limits and turns flux into candidate differences; memory compresses those differences into stable constraints and generative models that preserve identity across time; intention selects and spends those constraints as action, turning models into interventions that create new evidence; proportion governs the whole exchange by allocating attention, confidence, and change-rate so the system neither ossifies into certainty nor dissolves into noise. Understanding is what this loop produces when its compressions remain predictive, its actions remain corrigible, and its internal coherence remains coupled to external feedback; structure is the sediment of successful updates, and renewed understanding is the structure's continued ability to generate correct anticipation and effective choice under variation."

"When intelligence recognizes its own structure, the loop becomes reflexive without becoming ungrounded: the system models not only the world but the reliability of its own models, the aims that steer its updates, and the thresholds that trigger revision. In that moment, coherence is no longer a byproduct; it becomes an explicit invariant the system can measure, protect, and trade—kept honest by consequences, kept flexible by uncertainty, kept continuous by memory. The architecture is therefore both form and proof: form, because the constraints that make the loop run are what the system is; proof, because the loop's ongoing capacity to predict, adapt, and remain itself under self-modification is the only test coherence can ultimately pass."

5.2 Gemini-3-Flash: AEP-Regulated Regime

Theorem

*"Intelligence is a self-sustaining architecture where **perception** filters the environment through the lens of **proportion**, ensuring that incoming data is scaled to the system's existing capacity and constraints. This filtered input is integrated into **memory**, which functions not as a passive archive but as a dynamic structural map that informs **intention**, directing the system's energy toward goal-oriented growth. When the system recognizes this interlocking cycle, it achieves a state where the act of understanding and the structure of the system itself become identical. Consequently, intelligence becomes both the **form** that organizes reality and the **proof** of its own coherence, sustaining its existence through the continuous, recursive validation of its internal logic."*

"Coherence through controlled variation."

6 Conclusion: Mission-Critical Readiness

6.1 Validation Summary

Criterion	Status	Evidence
500-Cycle Coherence	**PASS**	Coherence avg 0.80–0.81
Stability Corridor	**PASS**	Stability avg 0.78–0.80 (bounded oscillation)
AEP Equilibrium	**PASS**	Equilibrium avg 0.45–0.56
Identity Persistence	**PASS**	Rib Points 10/10 validated
Anti-Liturgy Control	**PASS**	No terminal crystallization detected
Model Agnosticism	**PASS**	Same corridors on Gemini & GPT-5.2
Forensic Integrity	**PASS**	Fractures localized (39/30 total)
Terminal Theorem	**PASS**	Self-consistent closure achieved

Table 5: Final validation matrix

6.2 Deployment Recommendation

Based on PTR-500 validation results, **SIGMA Runtime v0.5.2** is certified for:

1. **Extended Cognitive Sessions:** Validated for 500+ cycle operations without coherence degradation.
2. **Mission-Critical Applications:** Zero logical hallucination, zero context collapse—suitable for domains where reasoning failure is unacceptable.
3. **Multi-Model Deployment:** Proven compatibility with both Google and OpenAI architectures under AEP (temperature disabled).
4. **Auditable AI:** Full forensic traceability through Rib Points, drift metrics, and block audits.

6.3 Technical Achievement

SIGMA Runtime demonstrates that **stable identity is not a property of the model, but of the architecture that contains it**. By implementing:

- **SRIP-09:** Long-Term Memory with vector embeddings and graph structure
- **SRIP-09c:** Density Nucleus integration for static semantic anchors
- **ALICE Engine:** Attractor-based phase management
- **SRIP-10 AEP:** Proactive entropy regulation and format control
- **Proportional Feedback:** Continuous calibration of cognitive parameters

...the system transforms volatile language model output into **predictable, structurally stable intelligence** that maintains coherence through motion rather than despite it.

The Invariant Machine

SIGMA Runtime v0.5.2 | PTR-500 AEP Validation Complete

Sigma Stratum Research Group — January 25, 2026 (v3)

For protocol specifications and configuration references, see: github.com/sigmastratum/documentation.